

Regular Expressions

CSC207 Fall 2015



Computer Science
UNIVERSITY OF TORONTO

The general idea

A **regular expression** is a pattern that a string may or may not match.

Example: `[0-9]+`

`[0-9]` means a character in that range

`“+”` means one or more of what came before

Strings that match: 9125 4

Strings that don't: abc empty string

Symbols like `[`, `,`, `]`, and `+` have special meaning. They are not part of the string that matches. (If we want them to be, we “escape” them with a backslash.)

Example Uses

Handling white space

A program ought to be able to treat any number of white space characters as a separator.

Identifying blank lines

Most people consider a line with just spaces on it to be blank.

Validating input

To check that input is has the expected format (e.g., a date in DD-MM-YYYY format).

Finding something within some input

E.g., finding dates within paragraphs of text.

Why bother?

We could accomplish those tasks without using patterns.

But it's much easier to declare a pattern that you want matched than to write code that matches it.

Therefore many languages offer support for this.

Bonus: By having the pattern explicitly declared, rather than implicit in code that matches it, it's much easier to:

- understand what the pattern is

- modify it

Editors, unix, Python, Java...

Regular expressions are used in many places.

Editors like vi, emacs, and Sublime Text allow you to use regular expressions for searching.

Many unix commands use regular expressions. Example:

```
grep pattern file
```

prints all lines from file that match pattern

Many programming languages provide a library for regular expressions.

The syntax varies from context to context, but the core is the same everywhere.

Simple patterns

Pattern	Matches	Explanation
<code>a*</code>	<code>''</code> <code>'a'</code> <code>'aa'</code>	zero or more
<code>b+</code>	<code>'b'</code> <code>'bb'</code>	one or more
<code>ab?c</code>	<code>'ac'</code> <code>'abc'</code>	zero or one
<code>[abc]</code>	<code>'a'</code> <code>'b'</code> <code>'c'</code>	one from a set
<code>[a-c]</code>	<code>'a'</code> <code>'b'</code> <code>'c'</code>	one from a range
<code>[abc]*</code>	<code>''</code> <code>'acbccb'</code>	combination

Note: In Java, patterns can be used to match an occurrence anywhere in the string, or one that consumes the whole string, among other options.

Anchoring

Lets you force the position of match

^ matches the beginning of the line

\$ matches the end

Neither consumes any characters.

Pattern	Text	Result
b+	abbc	Matches
^b+	abbc	Fails (no b at start)
^a*\$	aabaa	Fails (not all a's)

Escaping

Match actual

`^` and `$` and `[` etc.

using escape sequences

`\^` and `\$` and `\[` etc.

Remember, we also use escapes for other characters:

`\t` is a tab character

`\n` is a newline

Predefined Character Classes

Construct	Description
.	any character
\d	a digit [0-9]
\D	a non-digit [^0-9]
\s	a whitespace char [\t\n\x0B\f\r]
\S	a non-whitespace char [^\s]
\w	a word char [a-zA-Z_0-9]
\W	a non-word char [^\w]

The notation [^abc] means “anything not in the set”.

Defining your own Character Classes

Construct	Description
[abc]	a, b, or c (simple class)
[^abc]	any char except a, b, or c (negation)
[a-zA-Z]	a through z or A through Z inclusive (range)
[a-d[m-p]]	a through d or m through p (union)
[a-z&&[def]]	d, e, or f (intersection)
[a-z&&[^bc]]	a through z except for b and c (subtraction)
[a-z&&[^m-p]]	a through z and not m through p (subtraction)

Quantifiers

Construct	Description
$X?$	0 or 1 times
X^*	0 or more times
X^+	1 or more times
$X\{n\}$	exactly n times
$X\{n,\}$	at least n times
$X\{n,m\}$	at least n but no more than m times

Capturing Groups and Backreferences

Capturing groups allow you to treat multiple characters as a single unit.

Use **parentheses** to group.

Capturing groups are **numbered** by counting their opening parentheses from left to right.

((A)(B(C))) has the following groups:

1.((A)(B(C)))

2.(A)

3.(B(C))

4.(C)

Capturing Groups and Backreferences

The section of the input string matching the capturing group(s) is saved in memory for later recall via **backreference**.

A backreference is specified in the regular expression as a backslash (\) followed by a digit indicating the number of the group to be recalled.

For example,

Pattern	Example matching string
(\d\d)\1	1212
(\w*)\s\1	asdf asdf

Regular expressions in Java

The `java.util.regex` package contains:

`Pattern`: a compiled regular expression

`Matcher`: the result of a match

Example: `RegexDemo`